

PAVAN AKKINENI

☎ 6027567565 | ✉ pavanchowdary.akkineni@gmail.com | in [LinkedIn](#) | 📁 [Portfolio](#) | 🏠 [Databricks Cert](#) | 🌐 [AWS Cert](#)

EDUCATION

University of North Carolina at Chapel Hill

Masters in Computer Science, CGPA: 4.0/4.0

North Carolina, USA

Expected May 2026

- **Courses:** Deep Learning, LLMs for Software, Visual Computing, Image Processing, Multimodal AI, Reinforcement Learning and Markov Decision Processes, Machine Learning, Software Engineering.

WORK EXPERIENCE

UNC School Of Data Science and Society

Graduate Teaching Assistant

Chapel Hill, NC, USA

January 2025 – Present

- Mentored students in **statistics, machine learning, and neural networks**; designed and led hands-on workshops on clinical datasets; guided teams to build ML pipelines (ingestion, feature engineering, model selection, and evaluation).

UNC School of Medicine

Summer AI Engineering Intern

Chapel Hill, NC, USA

June 2025 - August 2025

- Engineered a predictive model for RNA MFE leveraging **Rinalmo(RNA LM)** context and K-mer pos-embeddings.
- Skilfully created antisense RNA-based therapeutic candidates for HIV TAR1 through **Evolutionary Algorithm**.

Bajaj Finance Limited

Senior Software Engineer

Pune, IN

September 2023 – August 2024

- Strategically steered **data-privacy** controls for **11+** databases, leveraging Securiti.ai, Databricks, and EDW(ASA) to achieve compliance with DPDP Act with RBAC/least-privilege IAM, PII masking, KMS-backed encryption.
- Engineered a high-availability Django platform (Securiti.ai-style) with async I/O, DB read-replicas, Celery queues, and Prometheus/Grafana observability—zero-downtime releases and $3.6\times$ faster scan cycles.
- Pioneered a data catalog/lineage service with REST/GraphQL contracts, RESTful microservices, schema versioning, and audit logging, enabling real-time metadata flow and idempotent writes securing \$80K/yr vendor savings.
- Automated an **OCR-driven** KYC ingestion pipeline using OpenCV, Tesseract with queue-backed workers, idempotent retries/backoff, PII redaction and structured logging; Graphana SLIs with alerts, enabling **60%** hands-free onboarding.
- Revamped legacy ingestion scripts to **Databricks Lakehouse** with **modular ETL**, schema enforcement **Delta**, **Unity Catalog ACLs** and lineage-aware audit logs, plus **AQE/shuffle tuning** across **24** production sources.

Data Engineer

July 2022 – August 2023

- Replatformed the entire production environment and scripts onto **Databricks**, reducing batch-upload time by 30%.
- Instituted expert-grade data-quality controls across OfferMart pipelines, eliminating **35%** of unnecessary processing.
- Overhauled scalable **Spark ETL** pipelines using **spark** and **Airflow** with lazy-loading, checkpointing and dynamic compute provision, cutting batch p95 by **25%** while orchestrating concurrent tasks for peak resource utilization.
- Led campaign offer generation with **DAG-based** policy workflows, queue-backed optimization and monitoring campaign dissonance through cron jobs with **SLA-backed** anomaly **alerts** and apply canary patches/rollback playbooks.

Data Engineer Intern

January 2022 – June 2022

- Contributed to Delta Lake data purging through VACUUM, retention policies, daily reports to save storage by **30%**.
- Facilitated partitions, z-order indexing strategies for **Hive** & Delta Lake tables, boosting p90 query latency by **25%**.

PROJECTS

- **MatchPlayer** - Masterminded a **YOLO-based** face-detection + **FaceNet-embedding** pipeline within an AI soccer commentary, driving player-ID accuracy to **52% Top-2 / 60% Top-3—about 15 × above the 4% random baseline**.
- **Prism LLM** - Expertly architected an **agentic** planner web app with **rate-limiting**, **retry/backoff** that classifies user query and dynamically routes it to the optimal LLM (DeepSeek, Mistral, Qwen, Gemma), boosting answer relevance.

TECHNICAL SKILLS

- **Data:** Spark SQL, Transact-SQL, PySpark, CDC (MERGE/Auto Loader), ADF, HDFS, MapReduce, DynamoDB, Azure Synapse Analytics, Azure Databricks, Azure Cost Management, ADF, Azure Devops, kafka, DynamoDB, MS Excel.
- **AI:** NLP, Computer Vision, Gen AI, LangChain, PyTorch, TensorFlow, Keras, scikit-learn, SciPy, NLTK, Hugging Face Transformers, MLflow, Data Acquisition and Handling
- **Other:** Python, C/C++, Java, Data Structures, OOP, Docker, Kubernetes, JavaScript, FastAPI, React, HTML, CSS, JQuery, Linux, Git, CRM, Cloud Computing, SDLC, CI/CD, SLIs/SLOs, Prometheus, Distributed tracing (OpenTelemetry), Agile/Scrum, Mentorship, Problem solving and Effective integration of related AI capabilities.

CERTIFICATIONS & AWARDS

- **Databricks certified Data Engineer Professional** -Structured streaming, Lakehouse, Testing, deployment, security.
- **AWS certified Machine Learning Associate** - Glue, S3, Kinesis, SageMaker, ECS, CloudWatch, Athena, Redshift.
- **Oracle AI Vector Search Certified Professional** - RAG, vector indexing, similarity search, Select AI, GoldenGate.
- **Bajaj Finserv Kudos Award** - Owning & leading the enterprise-wide high-stakes production replatforming to Databricks.